



give PRD

Below is a v0.1 PRD for the MVP: a privacy-first, no-community, full-screen discovery feed that makes replacing mindless scrolling feel effortless rather than educational or restrictive.

Product Requirements Document — “Useful Scroll” MVP

1. Product overview

Field	Definition
Working name	Useful Scroll — placeholder name; branding should not imply addiction or preachiness
Product type	Cross-platform mobile app for iOS and Android
Core platform	Flutter
Product stage	Serious startup experiment; functional MVP targeted in 4–6 weeks
Primary concept	A personalized vertical feed of short, trusted, useful content that can replace low-value social-media scrolling
Core promise	“Catch up on what matters and discover things you care about in a clean, private feed.”
Primary behavior to replace	Opening TikTok, Instagram, X, Reddit, or similar apps during boredom, anxiety, procrastination, commuting, waiting, breaks, or late-night wind-down
Core interaction	Full-screen vertical cards and short videos; swipe to continue; videos autoplay with sound
Community model	Zero community: no profiles, followers, comments, likes, DMs, uploads, or leaderboards
Privacy model	Local-first personalization; no account in V1; no server-side behavioral profile; no tracking-based advertising
Initial languages	English LTR and Hebrew RTL, including mixed-language rendering
Business direction	Freemium or optional local Pro features; no tracking-based ads

Problem statement

Existing short-form social feeds are extremely effective at filling idle moments, but they often leave users feeling drained, misinformed, distracted, or as though they wasted time. Existing news products can be overwhelming, politically loaded, slow, text-heavy, or stressful. Existing learning products usually require focus, commitment, explicit goals, and structured effort.

The opportunity is to create a feed-native experience that retains the low-friction, immersive quality of short-form scrolling while delivering concise, transparent, genuinely useful content:

current news, context, practical knowledge, hobbies, and personal interests.

Product thesis

A user should be able to open the app without planning what to learn, scroll naturally for a few minutes or longer, and leave knowing something relevant, interesting, or useful—without social pressure, creator drama, tracking-based ads, or emotionally manipulative content design.

Key tension

The product must not become “TikTok with better branding.” It must earn retention through relevance, trust, clarity, novelty, and user control—not through streaks, notifications, social validation, rage bait, or variable-reward manipulation.

The MVP must explicitly test whether users experience it as a meaningful replacement for traditional doom-scrolling rather than merely another feed.

2. Users and experience

Primary users

User	Trigger	Need	Product value
Tech-curious doom-scroller	Opens social media during breaks, procrastination, boredom, or commutes	Interesting, current, low-effort discovery	Personalized mix of AI, programming, homelab, self-hosting, science, markets, and global news
Overwhelmed news follower	Wants to know what happened without consuming hours of headlines	Clean, concise, non-sensational updates	Short multi-source news summaries, context, source transparency, and “why it matters”
Tired late-night user	Wants a low-energy distraction before sleep	Calm, non-graphic, lower-stimulation content	OLED dark mode, reduced-intensity feed selection, longer reading cards, no social pressure
Hobby-driven learner	Wants lightweight discovery in subjects they enjoy	Tips, stories, updates, inspiration, and deeper paths	Mixed content on music, guitars, scuba, gardening, gaming, sports, cars, and other interests
Privacy-sensitive user	Wants personalization without surveillance or community exposure	Local control and transparent data practices	No account in V1, local preferences, clear data deletion, no social graph

Core user journey

1. The user installs the app and chooses English or Hebrew.
2. The user selects broad interests during a lightweight onboarding flow.
3. The app immediately opens a personalized full-screen feed.
4. The user sees a concise news card, video preview, visual explainer, story, or hobby-related discovery item.

5. The user swipes vertically to continue.
6. The user can optionally save, share, request simpler wording, ask for more/less of a topic, inspect sources, or go deeper.
7. After sustained use, the app may show a gentle optional check-in or summary.
8. Saved content and recently viewed summaries remain available locally in the Library.

Product principles

1. **Useful by default.** Every item must provide a clear informational, practical, or discovery value.
2. **Trust before speed.** For factual content, especially news, transparent sourcing and uncertainty labels matter more than being first.
3. **Private by design.** Personal feed behavior and interest preferences remain local by default.
4. **No social pressure.** The product should not use public validation, popularity contests, social comparison, or community-driven engagement.
5. **Depth is optional.** The feed remains lightweight, but users can inspect sources and continue to original long-form content without losing their place.
6. **Control without clutter.** Users should shape the feed with simple actions rather than complex menus or mandatory configuration.
7. **Calm, not boring.** The app should be visually satisfying and fast while avoiding overstimulation, sensational visuals, and frantic interaction patterns.
8. **Transparent uncertainty.** The app must clearly distinguish confirmed facts, disputed claims, analysis, opinion, and developing stories.

In-scope feed formats

Format	Target length	Purpose	Go Deeper behavior
Breaking-news card	20–45 seconds	Explain what happened, what is confirmed, and why it matters	Sources, related timeline, original reporting
Daily news catch-up	60–120 seconds	Summarize major developments across selected categories	Open relevant original stories
Short video preview	20–60 seconds	Introduce a useful long-form video or creator content	Open in-app drawer/player or external source
Visual explainer	20–90 seconds	Explain a concept, chart, process, or comparison	Expanded explanation and sources
Practical tip card	20–45 seconds	Give a concise hobby, skill, or lifestyle insight	Source/original guide where relevant
Short story or discovery card	30–120 seconds	Deliver interesting context, history, science, culture, or niche-topic discovery	Expanded source or related items
“Why it matters” card	20–60 seconds	Provide context for market, geopolitical, technology, or societal developments	Source list and deeper explanation

Explicitly excluded from V1

- Social features of any kind.
- Comments, likes, follows, DMs, public profiles, public sharing feeds, or leaderboards.
- User uploads or user-generated posts.
- Standalone AI chat assistant.
- Quizzes, flashcards, streaks, XP, badges, levels, or gamified learning systems.
- Mandatory account creation or cloud sync.
- User-imported PDFs, RSS feeds, YouTube channels, newsletters, or web links.
- AI-generated videos from scratch.
- Keyword search.
- Tracking-based advertising.
- Full creator platform or creator monetization ecosystem.

3. Functional requirements

Onboarding and setup

ID	Requirement	Priority
ONB-01	On first launch, users select English or Hebrew.	Must
ONB-02	The interface must render fully in LTR and RTL, including mixed Hebrew/English text, metadata, buttons, dates, numbers, and source names.	Must
ONB-03	Users select broad interest categories during onboarding.	Must
ONB-04	Users may choose both news categories and personal-interest categories.	Must
ONB-05	The onboarding flow should remain lightweight and avoid asking users to self-assess knowledge level.	Must
ONB-06	The feed should open immediately after onboarding rather than landing on a dashboard.	Must
ONB-07	Users can modify their interests later in Settings.	Must
ONB-08	The product may use optional lightweight onboarding prompts such as “More tech or more world news?” to improve early feed quality.	Should

Feed behavior

ID	Requirement	Priority
FEE-01	The main screen opens directly to a full-screen vertical feed.	Must
FEE-02	Users move through the feed primarily by vertical swipe.	Must
FEE-03	The app supports mixed-topic sequencing in one unified feed.	Must
FEE-04	Every item displays a subtle topic tag, estimated consumption time, and source badge where relevant.	Must
FEE-05	Videos autoplay with sound when active in the viewport.	Must
FEE-06	The next video starts when the user swipes to it.	Must
FEE-07	The app must provide a user-accessible sound toggle and respect operating-system sound behavior.	Must
FEE-08	The first feed item and first swipe must load without visible stutter under normal network conditions.	Must
FEE-09	The feed should not include public popularity metrics such as views, likes, shares, or comments.	Must
FEE-10	Feed ranking must not depend on other users' behavior in V1.	Must
FEE-11	The app should avoid presenting repeated near-identical stories or the same topic too frequently.	Must
FEE-12	The app should use topic diversity rules so a user's feed does not become a narrow intellectual filter bubble.	Should
FEE-13	The app may include adjacent-topic exploration, such as cybersecurity for homelab users or music production for guitar users.	Should
FEE-14	The user can reduce or suppress politics, distressing news, or any unwanted topic through "Less like this."	Must

Item actions

ID	Requirement	Priority
ACT-01	Users can swipe away from any item with no interruption.	Must
ACT-02	Users can select "More like this."	Must
ACT-03	Users can select "Less like this."	Must

ID	Requirement	Priority
ACT-04	“Less like this” updates local preference weights without requiring a reason in V1.	Must
ACT-05	Users can select “Explain simply.”	Should
ACT-06	“Explain simply” must transform the current item into a simpler inline version, not open a chat experience.	Should
ACT-07	Users can select “Go Deeper” when an item has an eligible deeper source.	Must
ACT-08	Users can save an item to the local Library in one tap.	Must
ACT-09	Users can select “Sources” on every factual item.	Must
ACT-10	Users can share a clean summary snippet, source attribution, and item link through the operating system’s native share sheet.	Must
ACT-11	Share actions should support WhatsApp, Telegram, Messages, email, and other installed applications through the standard device share menu.	Must
ACT-12	The app must preserve the user’s feed position after saving, opening sources, sharing, or opening Go Deeper.	Must

Go Deeper behavior

ID	Requirement	Priority
DEP-01	Eligible content may provide a “Go Deeper” action.	Must
DEP-02	Go Deeper opens in an in-app drawer, sheet, or embedded preview where technically and legally permitted.	Must
DEP-03	The user can return to the exact prior feed position in one action.	Must
DEP-04	If embedded playback is unsupported or inappropriate, the app must offer an “Open original source” action.	Must
DEP-05	Original creators and publishers must receive clear attribution and direct source links.	Must
DEP-06	The product must not imply ownership of third-party content.	Must
DEP-07	V1 should initially prefer previews, metadata, authorized embeds, and source linking over technically complex or legally uncertain video clipping.	Must

Library, settings, and sessions

ID	Requirement	Priority
LIB-01	The app includes a Library tab.	Must
LIB-02	The Library displays saved items.	Must
LIB-03	The Library displays recently viewed summaries.	Should
LIB-04	The Library remains local to the device in V1.	Must
SET-01	Users can view and edit interest selections.	Must
SET-02	Users can select a desired session length.	Should
SET-03	The app may show a gentle optional check-in after extended use or after a selected session length.	Should
SET-04	The check-in must never shame, block, guilt, or force users to stop.	Must
SET-05	The app includes true-black OLED dark mode.	Must
SET-06	The app includes a night setting that can prioritize lower-stimulation and longer-form content.	Should
SET-07	Users can control whether distressing or conflict-related coverage appears, subject to the default non-graphic policy.	Must
SET-08	Users can clear all local data in one action.	Must
SET-09	Clear All Data removes local history, saved items, cached feed content, and local interest weights.	Must

4. Content, trust, and privacy

Content ingestion and normalization

The MVP content pipeline should ingest selected content from reputable RSS feeds, vetted news outlets, official and primary sources where available, specialist domain publications, and selected high-quality YouTube channels.

Every incoming item should be normalized into a common internal content schema:

Field	Description
Content ID	Unique internal identifier
Content type	News card, video preview, explainer, story, practical tip, or other approved format
Topic labels	One or more categories, such as AI, World News, Music, Markets, or Scuba
Language	Original language and rendered language

Field	Description
Headline	Concise, non-sensational title
Summary	Three to five factual bullets or a short visual narrative
Why it matters	Contextual explanation, clearly separated from confirmed facts
Source list	Original URLs, publisher/creator names, publication dates, and source tiers
Confidence status	Confirmed, developing, disputed, analysis, opinion, or low confidence
Risk classification	Standard, sensitive, high-risk, distressing, or review required
Go Deeper target	Original article, authorized video embed, source URL, or related context
Media asset	Permitted image, video preview, visual card, or thumbnail
Freshness	Time published, time summarized, and expiration/review window
Attribution	Required visible creator/publisher attribution

Editorial requirements

1. Every factual item must show a Sources button.
2. The source view must show the original links used to create the item.
3. High-risk content should use multiple independent sources whenever feasible.
4. Cards must clearly distinguish:
 - Confirmed facts.
 - Claims made by a named actor or source.
 - Analysis or interpretation.
 - Opinion.
 - Unverified or developing information.
5. If credible sources disagree, the card must identify the disagreement directly rather than presenting a false certainty.
6. The app must not use emotionally inflammatory headlines, rage-bait framing, or uncontextualized distressing imagery.
7. Major corrections must create an updated version of the content card and retain a visible correction note where appropriate.
8. Old news must not remain in the feed without contextual freshness labeling.
9. AI-generated summaries must never be treated as authoritative merely because the writing is fluent.
10. The system must avoid inventing causal explanations, motives, timelines, statistics, or implications absent from credible source material.

Sensitive-content policy

Topic	Minimum policy
Markets and investing	No buy/sell recommendations, price targets, personalized financial advice, certainty claims, or hype framing
Health and medicine	No diagnosis, treatment instruction, or personalized medical advice; require strong sourcing and caution labels
Politics and elections	Clear source diversity, distinction between reporting and opinion, and no false equivalence between evidence-based and unsupported claims
War and conflict	Non-graphic by default; clear uncertainty tags; avoid unverified casualty figures and sensational imagery
Law	No personalized legal advice; present context and primary-source links when feasible
Mental health	Avoid diagnosis or therapeutic claims; use careful language and point to professional resources if safety-relevant
Distressing events	Text-first summaries, context, user controls, and no graphic footage in default mode

Privacy requirements

ID	Requirement	Priority
PRV-01	No account or sign-in is required in V1.	Must
PRV-02	User preferences, saves, history, local feed interactions, and interest weights remain on-device by default.	Must
PRV-03	The backend must not maintain individual behavioral advertising profiles.	Must
PRV-04	The app must not sell user data or use tracking-based advertising.	Must
PRV-05	Personalization should run from local signals such as skips, saves, “more like this,” “less like this,” recency, and declared interests.	Must
PRV-06	V1 must not use aggregate cross-user recommendation intelligence to rank content.	Must
PRV-07	If anonymous telemetry is included, it must be minimal, documented, opt-out capable where practical, and limited to broad product-quality signals.	Must
PRV-08	Telemetry must not include raw personal reading history, saved-content lists, precise personal-interest profiles, or identifiable content-consumption trails.	Must
PRV-09	The app must provide an understandable privacy explanation in plain language.	Must
PRV-10	Clear All Data must wipe user-controlled local data immediately.	Must
PRV-11	Offline cache should preload roughly 20–30 cards without exposing local history to the backend.	Should

Safety assumptions to validate

The phrase “unbiased news” should not be implemented as a promise of total neutrality. The product requirement should instead be: transparent sourcing, neutral factual language, explicit uncertainty, meaningful viewpoint diversity, and visible editorial standards.

The MVP should avoid promising a fully autonomous “multi-source verified” global news engine before proving that the underlying source pipeline, story clustering, citation handling, and AI quality controls work reliably. Initial content scope should be limited to curated categories and a controlled source list.

5. MVP delivery and success

MVP scope

Must build	Simplify or operate manually	Defer
Flutter iOS/Android application	Curated initial source list	User source imports
English LTR and Hebrew RTL support	Manual editorial review for selected sensitive stories	Accounts and cloud sync
Lightweight interest onboarding	Limited topic taxonomy	Search
Full-screen vertical mixed feed	Start with text cards, source links, thumbnails, and authorized video previews	AI-generated video
Save, share, More Like This, Less Like This, Sources, and Go Deeper	Use an initial rules-based local ranking system	Advanced vector personalization
Local Library	Manual source-tier configuration	Cross-user collaborative filtering
Local preference storage	Limited daily content volume	Social features of any type
OLED dark mode	Simple optional session-length setting	Quizzes, streaks, levels, XP
Clear All Data	Basic offline cache	Full personalized learning assessment
RSS/news ingestion pipeline	Human review of flagged content	Broad long-tail global topic coverage
AI summary generation with strict source metadata	Controlled set of YouTube channels and embed/link flows	Automated video clipping at scale
Basic safety labels and source transparency	Limited analytics dashboard	Tracking-based ads

MVP acceptance criteria

The MVP is ready for a closed beta when:

- A user can install the Flutter app on iOS or Android and complete onboarding without an account.
- The app opens directly into a personalized full-screen feed.

- The feed supports at least several core categories, including general/world news, technology/AI, and selected hobby content.
- Content displays in English and Hebrew with correct bidirectional layout.
- Users can swipe smoothly between cards and video-preview items.
- Users can save an item, share it, inspect sources, choose More Like This or Less Like This, and return from Go Deeper without losing feed position.
- Every factual item displays a source path.
- News cards can display a “Developing,” “Disputed,” or equivalent uncertainty state.
- Sensitive-content defaults suppress graphic visual media.
- Local history, saved items, and preference weights can be deleted through Clear All Data.
- Offline cache can retain a limited next batch of feed items.
- The product has a basic event model for measuring anonymous activation, saving, source opens, feed completion, and session behavior without collecting a user-level server profile.

Success metrics

The product should not optimize for raw time spent alone. A user staying for 40 minutes can indicate value, but it can also indicate that the product recreated compulsive use under a healthier label.

Metric type	Metric	Why it matters
Activation	Percentage of new users who select interests and consume at least five items	Measures whether the initial experience is understandable
Feed relevance	Save rate, More Like This rate, Less Like This rate, and early skip rate by topic	Measures whether the feed is relevant without requiring a social graph
Trust	Source-button open rate, reported trust rating, correction rate, and user-reported clarity	Measures whether transparency is meaningful
Depth	Go Deeper usage, original-source opens, saved-item revisits, and voluntary return to a story	Measures whether the feed creates genuine curiosity
Retention	Day 1, Day 7, and Day 30 retention	Measures whether the product becomes a habit
Replacement	Self-reported reduction in traditional social-feed use	Measures the actual core promise
Satisfaction	“Did this feel like useful time?” post-session or periodic survey	Measures perceived value
Wellbeing	“Did you feel better, worse, or unchanged after this session?” and optional session-control usage	Detects whether the product is becoming another drain
Safety	Rate of user reports for inaccuracy, bias, distress, repetition, or low-quality content	Measures content quality and editorial risk
Operational	Cost per active user, average AI-processing cost per item, review burden, content freshness, and pipeline failure rate	Measures whether the product can scale sustainably

Initial success targets

These are directional validation targets rather than promises:

- Users can understand the feed and its controls without explanation.
- Users report that the content feels meaningfully more useful than a normal social-media feed.
- Users voluntarily return at least several times during a one-week beta.
- A meaningful portion of users save items, open sources, or use Go Deeper rather than only swipe endlessly.
- Users report that the product helps replace at least some social-media scrolling.
- Users do not describe the app as stressful, preachy, morally judgmental, manipulative, or indistinguishable from ordinary doom-scrolling.
- The content pipeline can reliably produce a small, high-quality daily feed from a controlled source set.

4–6 week plan

Period	Product and design	Engineering	Content and trust
Week 1	Define content-card templates, onboarding, primary feed interactions, library, and source view	Set up Flutter project, navigation, local storage, design system, LTR/RTL foundation	Select narrow initial source list; define content schema, source tiers, and sensitive-topic rules
Week 2	Build and test clickable feed prototype; refine swipe, actions, typography, dark mode	Implement feed, card rendering, onboarding, local preferences, save flow, Library	Build RSS ingestion proof of concept and manual curation workflow
Week 3	Design Go Deeper, source drawer, session check-in, and privacy settings	Build backend ingestion pipeline, content normalization, feed API, offline cache, source screen	Add AI summarization prompts, citation requirements, confidence labels, and manual review queue
Week 4	Conduct usability sessions; refine visual hierarchy and action discoverability	Integrate video-preview/link workflow, share sheet, telemetry, Clear All Data, Hebrew RTL QA	Create daily content operations process; test source conflicts, sensitive stories, corrections
Week 5	Closed beta polish; validate first-swipe performance and night mode	Device testing, feed performance, failure states, caching, analytics dashboard	Review card quality, source transparency, repetition, and political/news framing
Week 6	Analyze beta feedback; decide narrow next iteration	Fix critical issues and stabilize release candidate	Decide whether to expand topics, deepen personalization, improve video workflows, or narrow the product thesis

Pre-mortem

If the product fails after 12 months, the most likely reasons are:

1. It becomes another addictive feed, so users do not feel materially better after using it.
2. The content is summarized cleanly but feels generic, repetitive, or less entertaining than existing social platforms.
3. “Unbiased” becomes an unfulfillable brand promise and creates trust problems whenever users disagree with source selection or framing.
4. AI-generated summaries make factual errors, flatten important nuance, or create misleading confidence.
5. Video and creator-content workflows are too constrained by rights, platform rules, cost, or technical limitations.
6. A broad mix of news, AI, hobbies, markets, and lifestyle content creates a feed that feels unfocused rather than magical.
7. The local-first privacy model limits personalization quality or creates infrastructure complexity without being understood or valued by users.
8. The app lacks a sharp initial niche and cannot persuade users to replace entrenched habits.
9. Human editorial review becomes too expensive or too slow to maintain the trust standard.
10. Retention is weak because users prefer either entertainment feeds, specialist news products, or intentional learning apps rather than a hybrid.

The recommended next step is to use this PRD to create a narrower MVP design specification: first define the initial source list, first 5–8 interest categories, exact content-card templates, source-transparency rules, and the smallest viable content pipeline before writing the full Flutter feed.